Relay Protocol Specification

Section 1: Overview

The relay gateway accepts OpenAI compatible requests and queues them.
Workers consume jobs from Redis Streams with consumer groups and XAUTOCLAIM.
Rate limits use a Lua token bucket so that read and write stay atomic.

Section 2: Limits
The default request limit is 120 requests per minute.
Queue depth above 500 returns HTTP 503 with Retry-After.
The relay gateway accepts OpenAI compatible requests and queues them.
Workers consume jobs from Redis Streams with consumer groups and XAUTOCLAIM.
Rate limits use a Lua token bucket so that read and write stay atomic.

Section 3: Error codes

E401 means the API key is missing or revoked.

E429 means a rate limit was exceeded.

The relay gateway accepts OpenAI compatible requests and queues them.

Workers consume jobs from Redis Streams with consumer groups and XAUTOCLAIM.

Rate limits use a Lua token bucket so that read and write stay atomic.